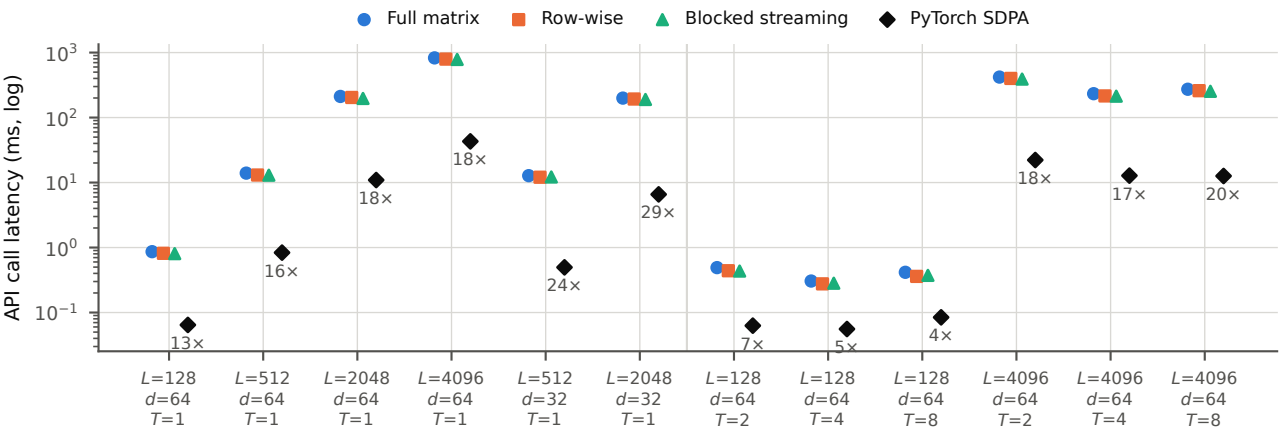
Full matrix
Row-wise
Blocked streaming
PyTorch SDPA
API call latency (ms, log)
10^3
10^2
10^1
10^0
10^{-1}
13×
16×
18×
18×
24×
29×
7×
5×
4×
18×
17×
20×
L=128 d=64 T=1
L=512 d=64 T=1
L=2048 d=64 T=1
L=4096 d=64 T=1
L=512 d=32 T=1
L=2048 d=32 T=1
L=128 d=64 T=2
L=128 d=64 T=4
L=128 d=64 T=8
L=4096 d=64 T=2
L=4096 d=64 T=4
L=4096 d=64 T=8