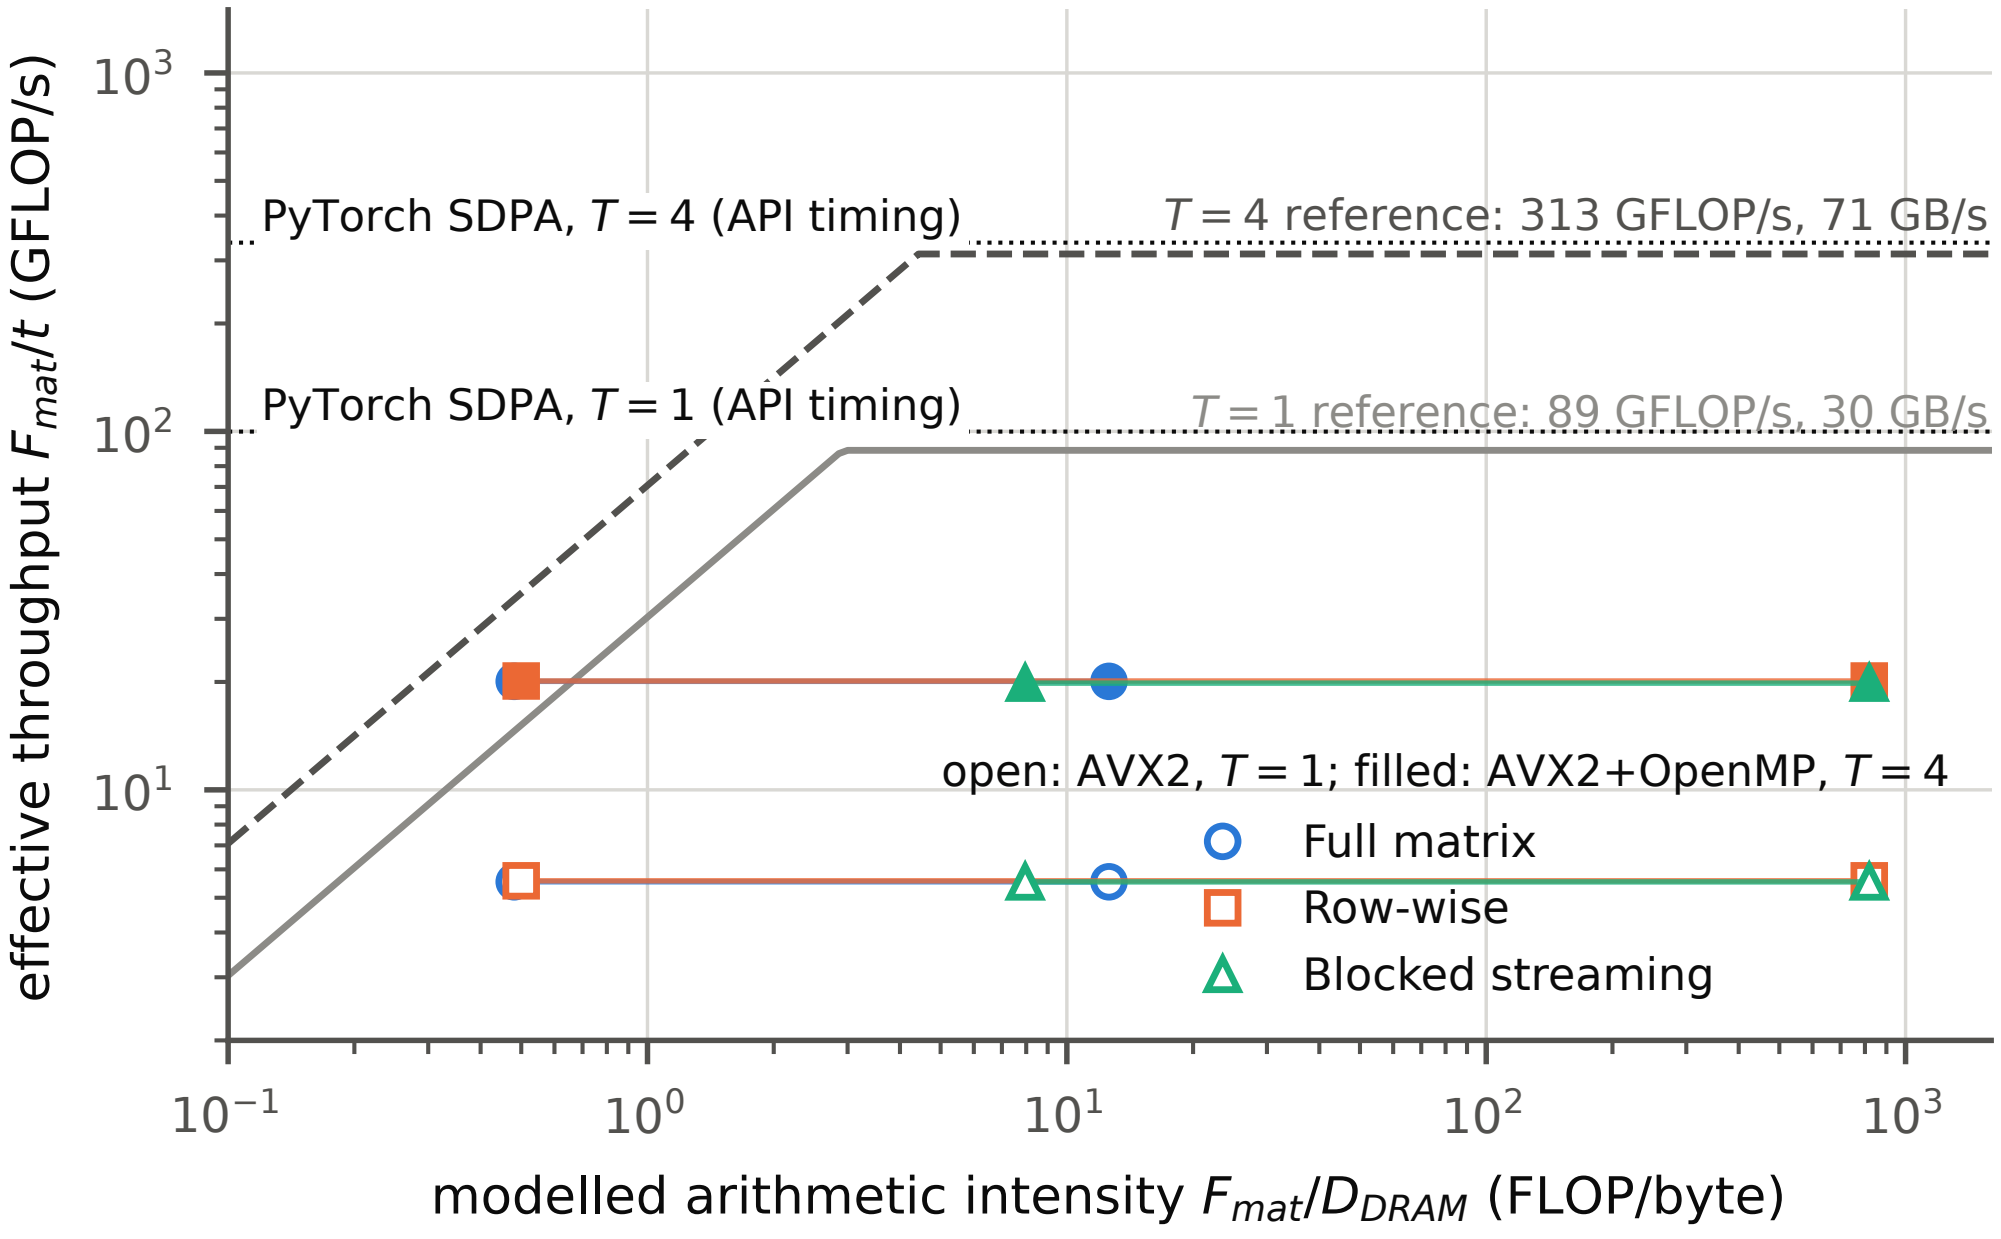
PyTorch SDPA, $T = 4$ (API timing)
$T = 4$ reference: 313 GFLOP/s, 71 GB/s
PyTorch SDPA, $T = 1$ (API timing)
$T = 1$ reference: 89 GFLOP/s, 30 GB/s
open: AVX2, $T = 1$; filled: AVX2+OpenMP, $T = 4$
Full matrix
Row-wise
Blocked streaming
effective throughput F_{mat}/t (GFLOP/s)
modelled arithmetic intensity F_{mat}/D_{DRAM} (FLOP/byte)
10^3
10^2
10^1
10^{-1}
10^0
10^1
10^2
10^3